Lecture 4 — Processes

Jeff Zarnett
jzarnett@uwaterloo.ca

Department of Electrical and Computer Engineering
University of Waterloo

November 14, 2015

Early computers did exactly one thing.

Or at least, exactly one thing at a time.

Now the OS supports multiple programs running concurrently.

To manage this complexity, we have the process.

A process is a program in execution.

1. The instructions and data.
2. The current state.
3. Any resources that are needed to execute.

The process is a very important concept in operating systems.

Practically everything the OS does involves one or more processes.

All other concepts we will examine in the future depend on this subject.

When we discuss scheduling, the problem is scheduling of processes.

The concept of process is fundamental to the structure of modern computer operating systems. Its evolution in analyzing problems of synchronization, deadlock, and scheduling in operating systems has been a major intellectual contribution of computer science.

- What Can Be Automated?: The Computer Science and Engineering Research Study, MIT Press, 1980

Most requirements of an operating system revolve around processes.

Some examples:

- **Scheduling**
- **Resource Allocation**
- **Inter-Process Communication**

Note: two instances of the same program running equals two processes.

You may have two windows open for Microsoft Word, and even though they are the same program, they are separate processes.

Similarly, two users who both use Firefox at the same time on a terminal server are interacting with two different processes.

Process Control Block

Data structure for managing processes: Process Control Block (PCB).

It contains everything the OS needs to know about the program.

It is created and updated by the OS for each running process.

It can be thrown away when the program has finished executing and cleaned everything up.

The blocks are held in memory and maintained in some container (e.g., a list) by the kernel.

The process control block will (usually) have:

- **Identifier.**
- **State.**
- **Priority.**
- **Program Counter.**
- **Register Data.**
- **Memory Pointers.**
- **I/O Status Information.**
- **Accounting Information.**

Identifier
State
Priority
Program counter
Memory pointers
Context data
I/O status information
Accounting information
• • •
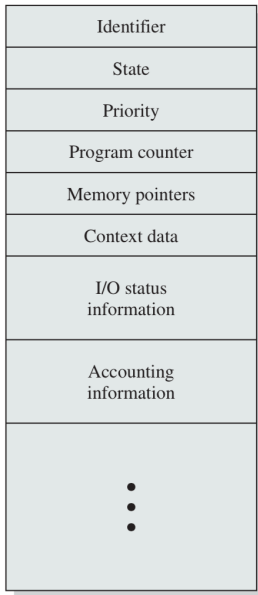
This data is kept up to date constantly as the process executes.

The program counter and the register data are asterisked.

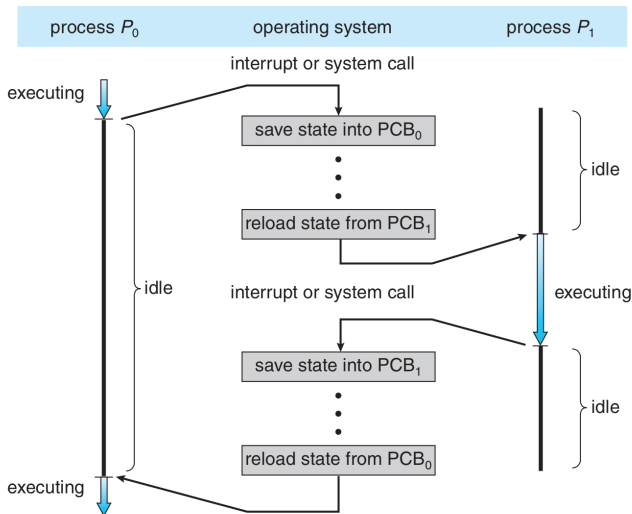When the program is running, these values do not need to be updated.

They are needed when a system call (trap) or process switch occurs.

We will need a way to restore the state of the program.

Save the state of the process into the PCB.

1. Save the state into the Program Counter variable.

2. Save the Register variables.

A switch between the execution of process P_0 and process P_1:

Unlike energy, processes may be created and destroyed.

Upon creation, the OS will create a new PCB for the process.
Also initialize the data in that block.

Set: variables to their initial values.
the initial program state.
the instruction pointer to the first instruction in `Main`

Add the PCB to the set.

After the program is terminated and cleaned up:
Collect some data (like a summary of accounting information).
Remove the PCB from its list of active processes and carry on.

Three main events that may lead to the creation of a process:

1 System boot up.
2 User request.
3 One process spawns another.

When the computer boots up, the OS starts and creates processes.

An embedded system might have all the processes it will ever run.

General-purpose operating systems: allow one (both) of the other ways.

Some processes will be in the foreground; some in the background.

A user-visible process: log in screen.

Background process: server that shares media on the local network.

UNIX term for a background process is Daemon.

Example: `ssh` (Secure Shell) command to log into a Linux system.

Users are well known for starting up processes whenever they feel like it.
Much to the chagrin of system designers everywhere.

Every time you double-click an icon or enter a command line command (like `ssh` above) that will result in the creation of a process.

An already-executing process may spawn another.

E-mail with a link? Click it; the e-mail program starts the web browser.

A program may break its work up into different logical parts.
  To promote parallelism or fault tolerance.

The spawning process is the parent and the one spawned is the child.

Eventually, most processes die.

This is sad, but it can happen in one of four ways:

1. Normal exit (voluntary)
2. Error exit (voluntary)
3. Fatal Error (involuntary)
4. Killed by another process (involuntary)

Most of the time, the process finishes because they are finished.
Or the user asks them to.

Compiler: when compilation is finished, it terminates normally.

You finish writing a document in a text editor, click the close button.

Both examples of normal, voluntary termination.

Sometimes there is voluntary exit, but with an error.

Required write access to the temporary directory & no permission.

Compiler: exit with an error if you ask it to compile a non-existent file.

The program has chosen to terminate because of the error.

Process Destruction: Fatality

The third reason for termination is a fatal error.
　Examples: stack overflow, or division by zero.

The OS will detect this error and send it to the program.

Often, this results in involuntary termination of the offending program.

A process may tell the OS it wishes to handle some kinds of errors.

If it can handle it, the process can continue.

Otherwise, unhandled exceptions result in involuntary termination.

The last reason for termination: one process might be killed by another.
  (Yes, processes can murder one another. Is no-one safe?!).

Typically this is a user request:
  a program is stuck or consuming too much CPU...
  the user opens task manager (Windows) or ps (UNIX)

Programs can, without user intervention, theoretically kill other processes.

Example: a parent process killing a child it believes to be stuck.

There are restrictions on killing process.
  A user or process must have the rights to execute the victim.

Typically a user may only kill a process he or she has created.

Exception: system administrator.

While killing processes may be fun, it is something that should be reserved for when it is needed.

Maybe when a process is killed, all processes it spawned are killed too.

In UNIX, but not in Windows, the relationship between the parent process and child process(es), if any, is maintained, forming a hierarchy.

A process, unlike most plants and animals, reproduces asexually.

A process has one parent; zero or more children.

A process and all its descendants form a process group.

Certain operations like sending a signal can be sent to a whole group.

UNIX the first process created is called `init`.

It is the parent of all processes (eventually).
 Like the `Object` class in Java is the superclass of all classes.

Thus in UNIX we may represent all processes as a tree structure.

In Windows, a parent process gets a reference to its child.

This allows it to exercise some measure of control over the child.

This reference may be given to another process (adoption).

No real hierarchy. A process in UNIX cannot disinherit a child.

When a process terminates, it does so with a return code.
Just as a function often returns a value.

On the command line or double clicking an icon, return value is ignored.

In UNIX, a parent can get the code that process returns.

Usually, a return value of zero indicates success.
Other values indicate an error of some sort.

Normally there is some sort of understanding between the parent and child processes about what a particular code means.

When a child process finishes, until the parent collects the return value, the child continues in a state of "undeath" we call a zombie.

This does not mean that the process then shuffles around the system attempting to eat the brains of other processes.

It just means that the process is dead but not gone.

There is still an entry in the PCB list.
   And the process holds on to its allocated resources.

Only after the return value is collected can it be cleaned up.

Usually, a child process's result is eagerly awaited by its parent.

The `wait` call collects the value right away.

This allows the child to be cleaned up (or, more grimly, "reaped").

If there is some delay for some reason, the process is considered a zombie until that value is collected.

If a parent dies before the child does, the child is called an orphan.

In UNIX any orphan is automatically adopted by the `init` process.
...making sure all processes have a good home.

By default, `init` will just `wait` on all its child processes
And do nothing with the return values.

A program can be intentionally orphaned: to run in the background.

This would be cruel, except that processes, as far as anyone knows, do not have feelings.